sewts

Weekly Review

09/03/21

- **Tasks**
- Custom Implementation : Cloth environment with SAC approach ⚠ *In progress*
- Prepare SAC Pipeline Explanation ✓
- **To-Do Items for Next Week**
- Finish custom Implementation : Cloth environment integrate to SAC approach
- Modify cloth initialization in simulation
- **To-Do Later**
- Define reward function and action for new use case
- Explore usage of intermediate testing on simulation before sim-to-real transfer
- Define use-case (for different type of towels (colour, texture, etc.) / one type)
- Check the no. of episodes needed, check computational requirements

Cloth Manipulation – RL Problem

02/03/21

Agent

State/Observation , Reward

Action

Environment (Mujoco)

Goal of RL Agent is to maximize expected reward

sewts

Cloth Manipulation using random policy

02/03/21

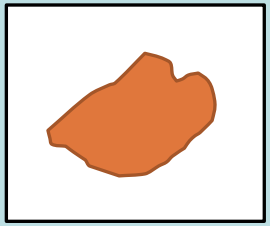

Observation
(64*64 img)

Reward
(Overlap with goal state)

Policy (state - action)=
Random action in a specific range

Agent

Action

Pick point and place point
From random pixel points
Inside segmented mask

Environment (Mujoco)

sewts

Cloth Manipulation using SAC

02/03/21

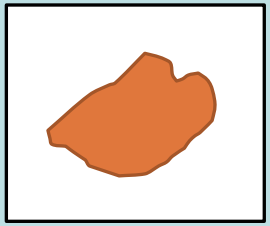

Observation
(64*64 img)

Reward
(Overlap with goal state)

SACAgent

Agent

Action

Pick point and place point

From random pixel points

Inside segmented mask

Environment (Mujoco)

STEP 1 : Sample actions from value network and store state, action, reward, next state, done in Experience Replay Buffer

Action sampled From mean, variance

Value Network (ψ)

Target Network (ψ)

Policy Network (φ)

Critic 1 Network (θ)

Critic 1 Network (θ)

Experience Replay Buffer

Maximize the reward + entropy of the policy

State, action, reward, next state, done

SAC Agent

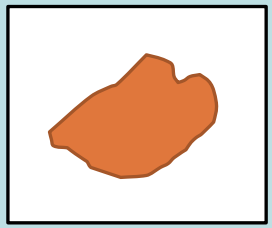

Observation (64*64 img)

Environment (Mujoco)

Action

Pick point

Place point

STEP 2 : We train the Quality Network by minimizing the following error

$$\hat{Q}(\mathbf{s}_t, \mathbf{a}_t) = r(\mathbf{s}_t, \mathbf{a}_t) + \gamma \mathbb{E}_{\mathbf{s}_{t+1} \sim p} \left[V_{\bar{\psi}}(\mathbf{s}_{t+1}) \right]$$

$$J_Q(\theta) = \mathbb{E}_{(\mathbf{s}_t, \mathbf{a}_t) \sim \mathcal{D}} \left[\frac{1}{2} \left(Q_\theta(\mathbf{s}_t, \mathbf{a}_t) - \hat{Q}(\mathbf{s}_t, \mathbf{a}_t) \right)^2 \right]$$

Value Network (ψ)

Target Network (ψ)

Policy Network (φ)

Critic 1 Network (θ)

Critic 1 Network (θ)

Experience Replay Buffer

Two networks are taken to avoid overestimation of Q values

Maximize the reward + entropy of the policy

State, action, reward, next state, done

SAC Agent

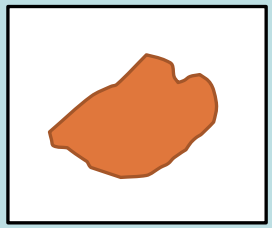

Observation (64*64 img)

Environment (Mujoco)

Action

Pick point

Place point

STEP 3 : We train the Value Network by minimizing the following error

$$J_V(\psi) = \mathbb{E}_{\mathbf{s}_t \sim \mathcal{D}} \left[\frac{1}{2} \left(V_\psi(\mathbf{s}_t) - \mathbb{E}_{\mathbf{a}_t \sim \pi_\phi} \left[Q_\theta(\mathbf{s}_t, \mathbf{a}_t) - \log \pi_\phi(\mathbf{a}_t | \mathbf{s}_t) \right] \right)^2 \right]$$

Value Network (ψ)

Target Network (ψ)

Policy Network (φ)

Critic 1 Network (θ)

Critic 1 Network (θ)

Experience Replay Buffer

State, action, reward, next state, done

Maximize the reward + entropy of the policy

SAC Agent

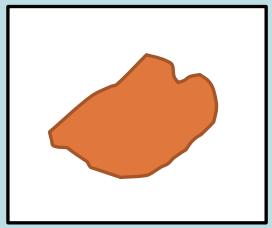

Observation (64*64 img)

Environment (Mujoco)

Action

Pick point

Place point

STEP 4 : We update the Target Network

Value Network (ψ)

Lags the Value Network

Target Network (ψ)

Policy Network (φ)

Critic 1 Network (θ)

Critic 1 Network (θ)

Target for Q Network depends on V Network and vice versa

Maximize the reward + entropy of the policy

Experience Replay Buffer

State, action, reward, next state, done

SAC Agent

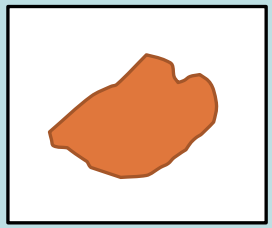

Observation (64*64 img)

Environment (Mujoco)

Action

Pick point

Place point

sewts

Plan

02/03/21

Planned

- Phase 3 : Implementation : 52 days (mid Feb- early Apr)

a) Testing various simulation environments and selecting one : 13 days
b) Setting up the Reinforcement Learning Platform and Simulation environment : 9 days
c) Dataset generation on chosen simulation platform : 15 days
d) Perform Reinforcement Learning using PyTorch : 15 days

Update

- Phase 3 : Implementation : 52 days (mid Feb- early Apr)

a) Setting up the Reinforcement Learning Platform and Simulation environment : 13 days
b) Prepare a custom implementation taking existing states, actions, rewards : 9 days
c) Redefine actions and rewards for our use case : 15 days
d) Test the pipeline and iterate : 15 days

Plan

02/03/21

- Phase 3 : Implementation : 52 days (mid Feb- early Apr)

a) Setting up the Reinforcement Learning Platform and Simulation environment : 13 days

b) Prepare a custom implementation taking existing states, actions, rewards : 9 days

c) Redefine actions and rewards for our use case : 15 days

d) Test the pipeline and iterate : 15 days

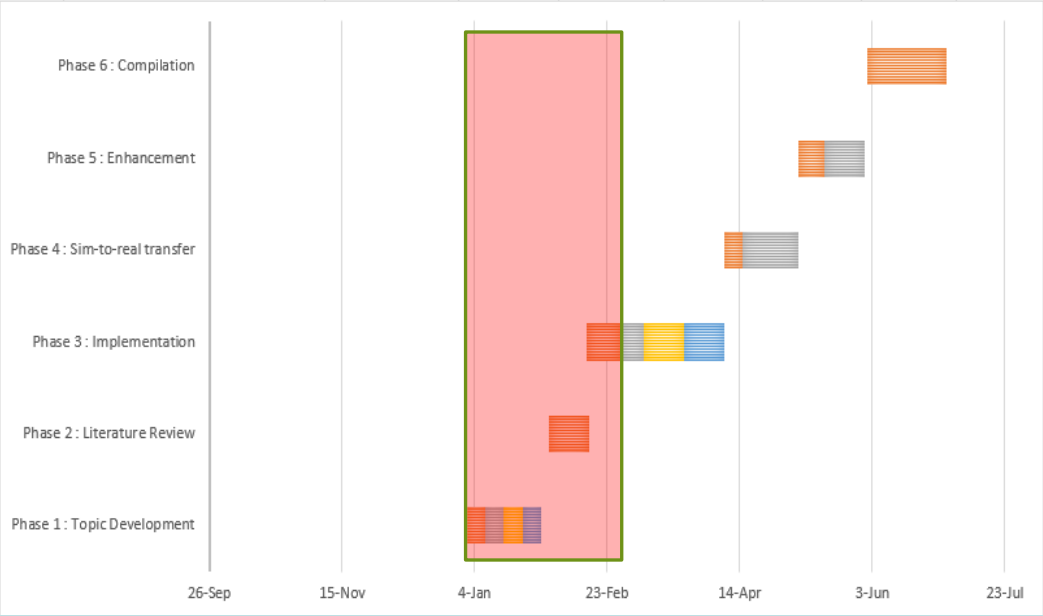

sewts

THANK YOU